

Horus Project Updates

Injection Attacks in LLMs

Ana Clara Zoppi Serpa

July 3

Past Presentation Recap

- Read [\[Gua+25\]](#)
- Cluster setup, model downloads, exploratory experiments inspired by [\[HNM26, Nou26\]](#)

Present

- Read paper [\[Rai+25\]](#) (helpful!)
- Found mech-interp collection on GitHub [\[Rai25\]](#), similar to the previous awesome-* series (awesome-jailbreak-papers, awesome-llm-papers etc.)
- nnsight tutorials [\[nns25\]](#)
- Research Questions exercise
- Found marimo [\[AS23\]](#)

References

- Agrawal, Akshay, and Myles Scolnick. 2023. “Marimo: An Open-Source Reactive Notebook for Python.”. <https://doi.org/10.5281/zenodo.12735329>.
- Guardieiro, Vitoria, Avishree Khare, Adam Stein, and Eric Wong. 2025. “Instruction Following by Principled Boosting Attention of Large Language Models.”. <https://arxiv.org/abs/2506.13734>.

Herring, Sam, Jake Naviasky, and Karan Malhotra. 2026. “Targeted Neuron Modulation via Contrastive Pair Search.”. <https://arxiv.org/abs/2605.12290>.

nnsight. 2025. “Nnsight Tutorials.”. <https://nnsight.net/tutorials/tutorials/>.

NousResearch. 2026. “Neural-Steering: Attribute and Steer Individual MLP Neurons in Language Models.”. <https://github.com/NousResearch/neural-steering>.

Rai, Daking. 2025. “Awesome Mechanistic Interpretability LM Papers.”. <https://github.com/Dakingrai/awesome-mechanistic-interpretability-lm-papers>.

Rai, Daking, Yilun Zhou, Shi Feng, Abulhair Saparov, and Ziyu Yao. 2025. “A Practical Review of Mechanistic Interpretability for Transformer-Based Language Models.”. <https://arxiv.org/abs/2407.02646>.